A N T H R O P I C

Claude AI Generalist

Complete Learning Journey

From Zero to Proficient — 8-Week Roadmap

PHASE 1

AI FOUNDATIONS

12 Topics • Deep Explanations • Exercises • Key Takeaways

Prepared By	**Mohan Kantipudi**
Prepared On	5th March 2026

Phase 1 — AI Foundations Overview

This phase builds the mental models every professional needs before using AI tools. You will understand exactly how AI thinks, why it sometimes fails, and how to work with it effectively.

What You Will Learn

Topic	What You Will Be Able To Do
1. What is AI	Explain AI to anyone in plain language
2. History of AI	Understand where AI came from and why it matters now
3. Machine Learning	Describe how machines learn from data
4. Large Language Models	Explain what LLMs are and how they are trained
5. How Claude & ChatGPT Work	Describe the architecture of modern AI chatbots
6. Tokens & Context Windows	Understand why AI has memory limits
7. Training vs Inference	Know the difference between building and using AI
8. Capabilities & Limitations	Set realistic expectations for AI tools
9. Hallucinations	Recognize and handle AI errors confidently
10. AI Safety & Ethics	Use AI responsibly and professionally
11. Responsible AI Usage	Apply ethical guidelines in daily work
12. AI Literacy	Position yourself as an AI-capable professional

1. What is Artificial Intelligence?

Simple Explanation

Artificial Intelligence (AI) is software that performs tasks which normally require human intelligence. This includes understanding language, recognizing images, making decisions, translating text, writing, coding, and much more.

Think of AI not as a robot with a brain, but as a very advanced calculator — one that has been trained on billions of examples of human knowledge and can predict useful responses.

The Analogy — Teaching a Child

ANALOGY	Imagine teaching a child to recognize cats. You show them thousands of pictures: 'cat', 'not cat'. After enough examples, the child can identify a cat they have never seen before. AI works exactly the same way — except instead of pictures, it trains on text, code, data, images, and more.

Visual Explanation

```
Traditional Software:        AI Software:

  INPUT                        INPUT
    ↓                            ↓
  RULES (written by humans)    MODEL (learned from data)
    ↓                            ↓
  OUTPUT                       OUTPUT

 You define every rule.      The model discovers patterns.
```

Three Types of AI You Should Know

Type	What It Does	Real Example
Narrow AI	Does one specific task very well	Chess engine, spam filter, face ID
General AI (AGI)	Can do any intellectual task (theoretical)	Does not yet exist in full form
Generative AI	Creates new content: text, images, code	Claude, ChatGPT, Midjourney, Gemini

Real-World Examples of AI Today

- Google Search ranking — AI decides which results to show you
- Netflix recommendations — AI predicts what you will enjoy
- Email spam filters — AI classifies messages automatically
- Voice assistants (Siri, Alexa) — AI understands spoken language
- Claude and ChatGPT — AI that can write, reason, and assist
- Medical diagnosis tools — AI detects diseases from scans

Beginner Exercise

Exercise 1: AI Spotter
Step 1: List 5 apps or services you use every day
Step 2: For each one, identify whether AI is involved Step 3: Write one sentence describing what the AI is doing in each case
Step 4: Discuss: Which felt most surprising?

Key Takeaways

1	AI is software that learns from data rather than following hard-coded rules.
2	Modern AI (like Claude) is called Generative AI — it creates new content.
3	AI is already in products you use daily without realizing it.
4	AI does not 'think' like humans — it predicts the most useful response based on patterns.

2. A Brief History of Artificial Intelligence

Why History Matters

Understanding where AI came from helps you appreciate why today's models are so powerful — and why they sometimes fail in predictable ways.

The AI Timeline

Era	Key Milestone
1950 — The Turing Test	Alan Turing asks: 'Can machines think?' He proposes testing machines by conversation.
1956 — The Word 'AI' is Born	John McCarthy coins the term 'Artificial Intelligence' at Dartmouth Conference.
1960s–70s — Early Programs	ELIZA chatbot created — one of the first programs to simulate conversation.
1980s — Expert Systems	AI encodes human expertise as rules. Limited but useful in specific domains.
1986 — Neural Networks Return	Backpropagation algorithm enables neural networks to learn from data.
1997 — Deep Blue	IBM's Deep Blue defeats chess world champion Garry Kasparov.
2012 — Deep Learning Revolution	AlexNet wins image recognition challenge. Deep learning dominates AI.
2017 — The Transformer	Google publishes 'Attention Is All You Need' — the paper that powers all modern LLMs.
2020 — GPT-3	OpenAI releases GPT-3. First model that can write convincingly like a human.
2022 — ChatGPT Launch	100 million users in 2 months. AI becomes mainstream for the first time.
2023 — Claude & the AI Race	Anthropic launches Claude. The AI assistant era begins in earnest.
2024–25 — AI Integration Everywhere	AI is embedded in every major product, workflow, and industry globally.

The Transformer — The Most Important Invention in AI

In 2017, Google researchers published a paper titled 'Attention Is All You Need.' This introduced the Transformer architecture, which is the foundation of every modern AI language model — Claude, GPT-4, Gemini, and others.

KEY INSIGHT	Before Transformers, AI models read text word by word, like reading a book one letter at a time. Transformers can read entire passages at once and understand relationships between all words simultaneously — like reading a whole page in one glance.

Key Takeaways

1	AI has been in development since the 1950s — it is not a new idea.
2	The 2017 Transformer paper is the single most important breakthrough in modern AI.
3	The AI explosion since 2022 is the result of decades of research finally converging.
4	We are currently living through the most significant AI transition in history.

3. What is Machine Learning?

Simple Explanation

Machine Learning (ML) is a subset of AI where systems learn from data without being explicitly programmed with rules. Instead of writing 'if it has whiskers and four legs, it is a cat,' you show the system thousands of cat photos and it figures out the pattern itself.

The Three Types of Machine Learning

Type	How It Learns	Example
Supervised Learning	Learns from labeled examples (input → correct answer)	Spam detection, image classification
Unsupervised Learning	Finds hidden patterns in unlabeled data	Customer segmentation, anomaly detection
Reinforcement Learning	Learns by trial and error with rewards and penalties	Game-playing AI, robot navigation

How Supervised Learning Works — Step by Step

```
STEP 1: Collect Data
  10,000 emails labeled: [spam] or [not spam]

STEP 2: Train the Model
  Show model email → model guesses → compare to label → adjust

STEP 3: Repeat Millions of Times
  Model gradually improves — errors get smaller

STEP 4: Deploy
  Model can now classify NEW emails it has never seen
```

ANALOGY	Machine learning is like teaching someone to ride a bike by letting them fall and recover, rather than reading a manual. The system learns from experience, not instructions.

Why Machine Learning Matters for AI Users

When you understand ML, you understand why AI models improve over time, why they need massive amounts of training data, and why the quality of training data directly affects the quality of AI outputs.

Key Takeaways

1	ML is AI that learns from data instead of fixed rules.
2	The most common type for language AI is supervised learning.
3	More high-quality data generally means better AI performance.
4	ML models do not 'understand' — they recognize and reproduce statistical patterns.

4. What Are Large Language Models?

Simple Explanation

A Large Language Model (LLM) is a type of AI trained on massive amounts of text. The model learns the statistical patterns of language — which words follow which, how ideas connect, what makes a good answer — and uses this to generate coherent, useful text.

The word 'large' refers to the number of parameters in the model. Modern LLMs have billions or even trillions of parameters — internal numerical values that the model adjusts during training.

The Anatomy of an LLM

```
TRAINING DATA          MODEL ARCHITECTURE         OUTPUT

Books ──────┐
Websites ───┤
Code ───────┤   →  [Transformer Neural Network] →  Predictions
Wikipedia ──┤          (Billions of parameters)    (Text responses)
Articles ───┘

SIZE of training data × QUALITY of training data × ARCHITECTURE = Model capability
```

Parameters — What Are They?

Parameters are the numbers inside a neural network that get tuned during training. Think of them like the knobs on a complex mixing board — millions to trillions of tiny adjustments that collectively allow the model to understand and generate language.

Model	Approximate Parameters
GPT-2 (2019)	~1.5 billion
GPT-3 (2020)	~175 billion
Claude 3 Haiku	Not disclosed, optimized for speed
Claude 3.5 Sonnet	Not disclosed, balanced performance
GPT-4 (estimated)	~1 trillion (unofficial estimate)

ANALOGY

If parameters are knobs, training is the process of automatically turning all those knobs until the model's outputs match human-quality text. A model with 100 billion parameters has 100 billion knobs being tuned simultaneously.

What LLMs Are Good At

- Writing — essays, emails, reports, stories, code
- Summarizing — condensing long documents into key points
- Translating — converting between languages
- Answering questions — drawing on training knowledge
- Following instructions — step-by-step task completion
- Reasoning — working through logic problems

What LLMs Are NOT

- Not a search engine — they do not browse the internet (unless a tool is provided)
- Not a database — they do not store or retrieve facts reliably
- Not sentient — they have no feelings, opinions, or consciousness
- Not infallible — they make mistakes, especially with recent events or precise numbers

Key Takeaways

1	LLMs are neural networks trained on vast amounts of text to predict useful language.
2	They are characterized by the number of parameters (billions to trillions).
3	LLMs generate responses based on pattern recognition, not true comprehension.
4	Claude, GPT-4, and Gemini are all examples of LLMs.

5. How Claude and ChatGPT Work

The Full Pipeline

When you type a message to Claude, a sophisticated multi-step pipeline processes it before you see a response. Understanding this helps you write better prompts and set accurate expectations.

```
YOU TYPE:      'Explain quantum computing simply'
      ↓
TOKENIZATION:  Your text is split into tokens (word pieces)
      ↓
ENCODING:      Tokens are converted to numbers (vectors)
      ↓
ATTENTION:     Model analyzes relationships between all tokens
      ↓
PREDICTION:    Model predicts the best next token, then the next...
      ↓
DECODING:      Numbers are converted back to words
      ↓
YOU SEE:      A clear, structured explanation of quantum computing
```

The Role of Reinforcement Learning from Human Feedback (RLHF)

Raw LLMs produce outputs that are statistically plausible but often unhelpful or unsafe. RLHF is the process that transforms a raw LLM into a helpful assistant like Claude.

Stage	What Happens
1. Pre-training	Model reads trillions of tokens and learns general language patterns
2. Supervised Fine-tuning	Human trainers write ideal responses; model learns to mimic them
3. Reward Model Training	Humans rank AI outputs; a reward model learns what 'good' looks like
4. RLHF	The LLM is trained using reinforcement learning to maximize the reward score
5. Constitutional AI (Anthropic)	Claude is also trained with written principles — making it safer and more aligned

Claude vs ChatGPT — How Are They Different?

Feature	Claude (Anthropic)	ChatGPT (OpenAI)
Safety Approach	Constitutional AI — trained on written principles	RLHF + human feedback

Context Window	Very large (200K tokens on Claude 3)	128K tokens (GPT-4o)
Personality	Thoughtful, nuanced, careful	Helpful, conversational, flexible
API Access	Yes — via api.anthropic.com	Yes — via api.openai.com
Document Handling	Excellent — can analyze large files	Good — similar capabilities

Key Takeaways

1	Claude processes your text through tokenization, attention, and prediction.
2	RLHF is what makes a raw LLM into a safe, helpful assistant.
3	Anthropic uses Constitutional AI — a unique approach that bakes in principles.
4	Claude and ChatGPT have similar foundations but different training philosophies.

6. Tokens and Context Windows

What is a Token?

AI models do not read word-by-word or letter-by-letter. They read in chunks called tokens. A token is typically 3-4 characters of text, or about 0.75 words on average.

```
Example sentence:
  'Artificial intelligence is transforming the world.'

Tokenized approximately as:
  ['Art', 'ific', 'ial', ' intel', 'ligence', ' is', ' transform', 'ing', ' the',
' world', '.']

Token count: ~11 tokens for this sentence
Rule of thumb: 1,000 tokens ≈ 750 words ≈ about 1.5 pages
```

Why Tokens Matter for You

Concern	Why It Matters
Cost	API pricing is per token — longer conversations cost more
Speed	More tokens = slower response generation
Memory	Models only 'remember' what fits in the context window
Quality	Very long prompts can dilute the model's focus

What is a Context Window?

The context window is the total amount of text an AI model can 'see' at once — including your message, the conversation history, any documents you uploaded, and its own previous responses.

ANALOGY	Think of the context window like a whiteboard in a meeting room. Everything written on the whiteboard is visible to everyone. When the whiteboard fills up, older information must be erased to make room. The AI can only work with what is currently on the whiteboard.

```
CONTEXT WINDOW (example: 200,000 tokens for Claude 3)

  ┌──────────────────────────────────────────────────────────┐
  │ System Prompt      │ Conversation History │ Your      │
  │ (instructions)     │ (all previous turns) │ Message   │
```

```
  └──────────────────────────────────────────────────────────────┘
   ←──────────────────── 200,000 tokens total ────────────────────→

200,000 tokens ≈ 150,000 words ≈ a full novel
```

Token Counting Reference

Content	Approximate Token Count
One tweet (280 chars)	~70 tokens
One email (200 words)	~270 tokens
One page of text (500 words)	~667 tokens
A 10-page report	~6,700 tokens
A full book (80,000 words)	~107,000 tokens
Claude 3 context window	200,000 tokens (about 1.5 books)

Key Takeaways

1	Text is split into tokens — about 750 words per 1,000 tokens.
2	The context window is the AI's working memory — it cannot access anything outside it.
3	Claude 3 has one of the largest context windows available — up to 200K tokens.
4	For very long documents, break them into sections or use retrieval methods.

7. Training vs Inference

The Two Phases of AI

Every AI model goes through two completely different phases: training (where it learns) and inference (where it works). As an AI user, you are always interacting during the inference phase.

	Training	Inference
What happens	Model learns from data	Model generates responses
Who does it	AI researchers at Anthropic	You (the user)
When it happens	Before the model is released	Every time you chat
How long it takes	Weeks to months	Milliseconds to seconds
Cost	Extremely expensive (millions $)	Much cheaper per call
Changes the model?	YES — adjusts parameters	NO — model stays the same

KEY INSIGHT

When you chat with Claude, you are NOT training it. Your conversation does not change Claude's underlying model. Each conversation starts fresh (unless memory features are enabled). You are using the model, not teaching it.

Why This Matters for Users

- Claude's knowledge has a cutoff date — it was trained on data up to a certain point
- Telling Claude 'that was wrong' in a chat does not permanently fix the model
- You cannot 'corrupt' Claude by asking it questions — inference is read-only
- Fine-tuning is when companies re-train a model on their own data — that is a separate process

Knowledge Cutoff

Because training happens on historical data, Claude's knowledge has a cutoff — a date after which it does not know about world events. Always check Claude's knowledge cutoff for time-sensitive topics.

Exercise 2: Training vs Inference Reflection

Step 1: Write down 2 questions you could ask Claude that it could answer from training data

Step 2: Write down 2 questions that require real-time information Claude would NOT know

Step 3: For the real-time questions, what tool or approach could you use instead?
Step 4: Reflection: Why does this distinction matter in professional work?

Key Takeaways

1	Training is when the model learns — inference is when you use it.
2	Chatting with Claude does not change the model permanently.
3	Claude has a knowledge cutoff date — it does not know recent events unless told.
4	For real-time information, combine Claude with web search or provide the information yourself.

8. AI Capabilities vs Limitations

Setting Realistic Expectations

One of the most important skills for an AI professional is knowing when to use AI and when NOT to. Mismatched expectations are the #1 cause of frustration with AI tools.

What Claude Excels At

Capability	Example Task
Text generation	Write a 2,000-word report on climate change policy
Summarization	Condense a 50-page legal document to key points
Code writing	Write a Python function to parse JSON data
Editing & proofreading	Fix grammar, improve clarity, restructure paragraphs
Brainstorming	Generate 20 marketing slogans for a coffee brand
Explanation	Explain blockchain to a 10-year-old
Translation	Translate a business email from English to French
Analysis	Analyze the strengths and weaknesses of this business plan
Structured tasks	Create a project plan with milestones and deadlines

What Claude Struggles With

Limitation	Why It Happens
Real-time information	Training data has a cutoff — no live internet access by default
Precise arithmetic	LLMs process tokens, not numbers — calculators are better for math
Long-term memory	Each conversation starts fresh unless memory tools are used
Factual recall of obscure facts	Training data may be sparse on niche topics
Physical world tasks	Claude generates text — it cannot click, buy, or act in the world
Consistent long documents	Very long outputs can lose coherence toward the end

The Golden Rule of AI Usage

RULE	Use AI as your first draft, not your final product. Claude accelerates your work dramatically, but always apply your judgment, expertise, and verification before publishing or acting on AI outputs.

The AI Confidence Problem

AI models speak confidently even when they are wrong. Unlike a human expert who might say 'I'm not sure,' Claude may state incorrect information in the same confident tone it uses for correct information. This is called the confidence calibration problem.

- Always verify critical facts with authoritative sources
- Ask Claude to note its uncertainty: 'Tell me what you know and flag anything uncertain'
- Cross-check important data — especially statistics, dates, and names

Key Takeaways

1	Claude excels at language tasks: writing, summarizing, explaining, coding, and analysis.
2	Claude has real limitations: no real-time data, imprecise math, no persistent memory.
3	The biggest risk is over-trusting confident-sounding but incorrect outputs.
4	Use Claude as a powerful collaborator, not an infallible oracle.

9. Understanding AI Hallucinations

What is a Hallucination?

An AI hallucination is when the model generates content that sounds plausible and confident but is factually incorrect. The term comes from psychology — producing false perceptions that feel real.

This is not lying. The AI has no intent. It is a natural consequence of how language models work: they are optimized to produce coherent, fluent text, not to produce only verified facts.

Examples of Hallucinations

Type	Example
Fake citations	'According to Smith et al. (2019)...' — paper does not exist
Wrong dates	Claims an event happened in the wrong year
Invented people	Describes a researcher or expert who does not exist
False statistics	States a percentage that sounds accurate but is fabricated
Wrong technical details	Describes how a protocol works incorrectly
Mixed-up entities	Confuses two similar companies, people, or places

Why Hallucinations Happen

```
LLMs predict the MOST PROBABLE next token, not the MOST TRUE next token.

Query: 'What papers did Professor X publish on topic Y?'
                              ↓
Model thinks: 'What text typically follows a question about a professor's papers?'
                              ↓
Model generates: A plausible-looking citation in the correct format
                              ↓
Reality: The citation may be entirely invented
```

How to Reduce Hallucinations

- Provide source material — paste the document, don't ask from memory
- Ask for uncertainty flags: 'If you are not sure, say so explicitly'
- Request citations with links — then verify them independently
- Use Claude for drafting and analysis, not as a fact database
- Break questions into smaller parts — complex multi-part queries hallucinate more

- Use web search + Claude together for fact-sensitive tasks

PRO TIP	The best way to minimize hallucinations is to give Claude the information and ask it to work with that, rather than asking it to recall facts from training. Think: 'Analyze this document' rather than 'Tell me about this topic from memory.'

Exercise 3: Hallucination Detection
Step 1: Ask Claude: 'Name 5 academic papers published about AI safety in 2022 with authors and journals'
Step 2: Pick 2 of the papers and search for them online
Step 3: Document whether they are real, partially real, or invented
Step 4: Reflection: What does this tell you about using AI for research?

Key Takeaways

1	Hallucinations are false outputs that sound confident and correct.
2	They happen because LLMs optimize for fluency, not factual accuracy.
3	The best defense is providing source material rather than asking from memory.
4	Always verify critical facts — especially citations, statistics, and names.

10. AI Safety and Ethics

Why Safety and Ethics Matter

AI is one of the most powerful technologies ever created. Used well, it can accelerate medicine, education, and productivity. Used carelessly, it can spread misinformation, amplify bias, violate privacy, and cause harm. As an AI professional, you have a responsibility to use these tools thoughtfully.

Key Ethical Concerns in AI

Concern	Description
Bias	AI models inherit biases from training data — they can reflect and amplify human prejudices
Privacy	AI processes vast personal data — consent and security are critical
Misinformation	AI can generate convincing but false content at scale
Job displacement	Automation may eliminate certain roles while creating others
Deepfakes	AI-generated media can deceive at scale
Surveillance	AI face recognition enables mass surveillance
Concentration of power	A few companies control the most powerful AI systems
Existential risk	Advanced AI could pose systemic risks if not aligned with human values

Anthropic's Approach — Constitutional AI

Anthropic was founded with AI safety as its core mission. Claude is trained using a technique called Constitutional AI (CAI), which means:

- Claude is trained with a written set of principles — a 'constitution'
- The model learns to critique and revise its own outputs against these principles
- This makes Claude safer without requiring constant human review of every output

Key Takeaways

1	AI ethics covers bias, privacy, misinformation, displacement, and systemic risk.
2	Anthropic was founded specifically to build AI safely — Claude reflects this mission.

3	Constitutional AI is a unique approach that trains Claude to follow written principles.
4	AI safety is not just for researchers — every user has a responsibility to use AI thoughtfully.

11. Responsible AI Usage

The Responsible AI Checklist

Use this checklist every time you deploy AI in a professional context:

Check	Question to Ask
Accuracy	Have I verified the AI's factual claims against authoritative sources?
Attribution	Is it clear where this content came from (human vs AI-assisted)?
Bias	Could this output contain biased language or unfair representations?
Privacy	Am I sharing personal or confidential data with the AI inappropriately?
Harm	Could this output cause harm to anyone if acted upon?
Transparency	Am I being honest with stakeholders about AI involvement?

What NOT to Put in AI Prompts

- Personal identifiers — names, social security numbers, ID numbers
- Private health information — diagnoses, medications, records
- Confidential business data — trade secrets, unreleased financials
- Passwords or API keys
- Third-party personal data without consent

RULE	Treat every AI prompt as a public document. If you would not be comfortable with the information becoming visible to others, do not include it in an AI prompt.

Key Takeaways

1	Always verify AI outputs before using them professionally.
2	Never share sensitive personal, health, or confidential business data in prompts.
3	Be transparent with colleagues and stakeholders about AI involvement in your work.
4	Apply the responsible AI checklist to every significant AI-assisted task.

12. AI Literacy for Professionals

What is AI Literacy?

AI literacy is the ability to understand, evaluate, and effectively use AI tools in your professional context. It is rapidly becoming as essential as computer literacy was in the 1990s.

The Four Levels of AI Literacy

Level	What It Means
Level 1 — Awareness	You know AI exists and can name major tools
Level 2 — User	You can use AI tools to complete tasks effectively
Level 3 — Collaborator	You design workflows that integrate AI as a team member
Level 4 — Builder	You create AI-powered applications using APIs and code

This course takes you from Level 1 to Level 4.

AI Skills Employers Are Looking For

- Prompt engineering — ability to extract high-quality outputs from AI
- AI tool proficiency — familiarity with Claude, ChatGPT, Copilot, Gemini
- Critical evaluation — ability to assess and verify AI outputs
- Workflow design — integrating AI into professional processes
- API usage — building with AI programmatically
- Ethical judgment — using AI responsibly and transparently

Exercise 4: AI Literacy Self-Assessment
Step 1: Rate yourself 1-5 on each of the 6 AI skills listed above (1 = no knowledge, 5 = expert)
Step 2: Identify your lowest-rated skill — that is your highest-priority area for development Step 3: Write a 3-sentence goal statement: What you want to achieve with AI in 8 weeks
Step 4: Revisit this assessment after completing all 8 phases of this course

Key Takeaways

1	AI literacy is a career-defining skill for every profession in the next decade.
2	There are four levels: awareness, user, collaborator, and builder.
3	This course is designed to take you to Level 4 — builder.
4	The best investment you can make right now is consistent daily practice with AI tools.

PHASE 1 — COMPLETE SUMMARY

Phase 1 Complete — What You Have Learned

Master Concepts

Concept	One-Line Summary
Artificial Intelligence	Software that learns from data to perform tasks requiring human-like intelligence
Machine Learning	AI that improves through exposure to training data, not hard-coded rules
Large Language Models	Neural networks with billions of parameters trained on vast text corpora
Transformers	The 2017 architecture that powers all modern AI language models
RLHF	Training technique that shapes raw LLMs into safe, helpful assistants
Constitutional AI	Anthropic's approach: training Claude with written ethical principles
Tokens	The chunks AI reads — ~750 words per 1,000 tokens
Context Window	The AI's working memory — what it can 'see' in one conversation
Training	When the model learns — happens before deployment, very expensive
Inference	When you use the model — happens every time you chat
Hallucination	Confident but false AI outputs — a natural consequence of pattern prediction
AI Literacy	The skill of effectively, ethically, and critically using AI in professional work

Phase 1 Assessment — Test Yourself

1. Without looking, explain what a token is in one sentence.
2. What is the difference between training and inference?
3. Name three things Claude is excellent at and three things it struggles with.
4. What is a hallucination and how can you reduce it?
5. What is Constitutional AI and why did Anthropic create it?
6. What is the context window and why does it matter for long documents?

Coming Up — Phase 2

NEXT	Phase 2: AI Collaboration & Prompt Engineering — Learn to write prompts that consistently produce professional-quality outputs. Master the techniques used by AI power users worldwide.